

AI FOR MATERIALS INDUSTRY

ARTIFICIAL INTELLIGENCE

A MASSIVE OPEN ONLINE COURSE

Hands-on session 2

Steel defect classification with deep learning




What is AI?



Series of techniques to extract **information** from data

Types of AI



The more data we have, the more we trade insight for predictive power

Case studies

Spreadsheet data




Steel




Computer vision

Glass



Materials discovery


Manufacturing




Sensor data

Steel plate defects


Crazing




Inclusions




Patches




Rolled-in scale




Scratches




Pits




Multi-step pipeline



Detector




Featurizer



Classifier

Can we make a model that classifies the raw image directly?

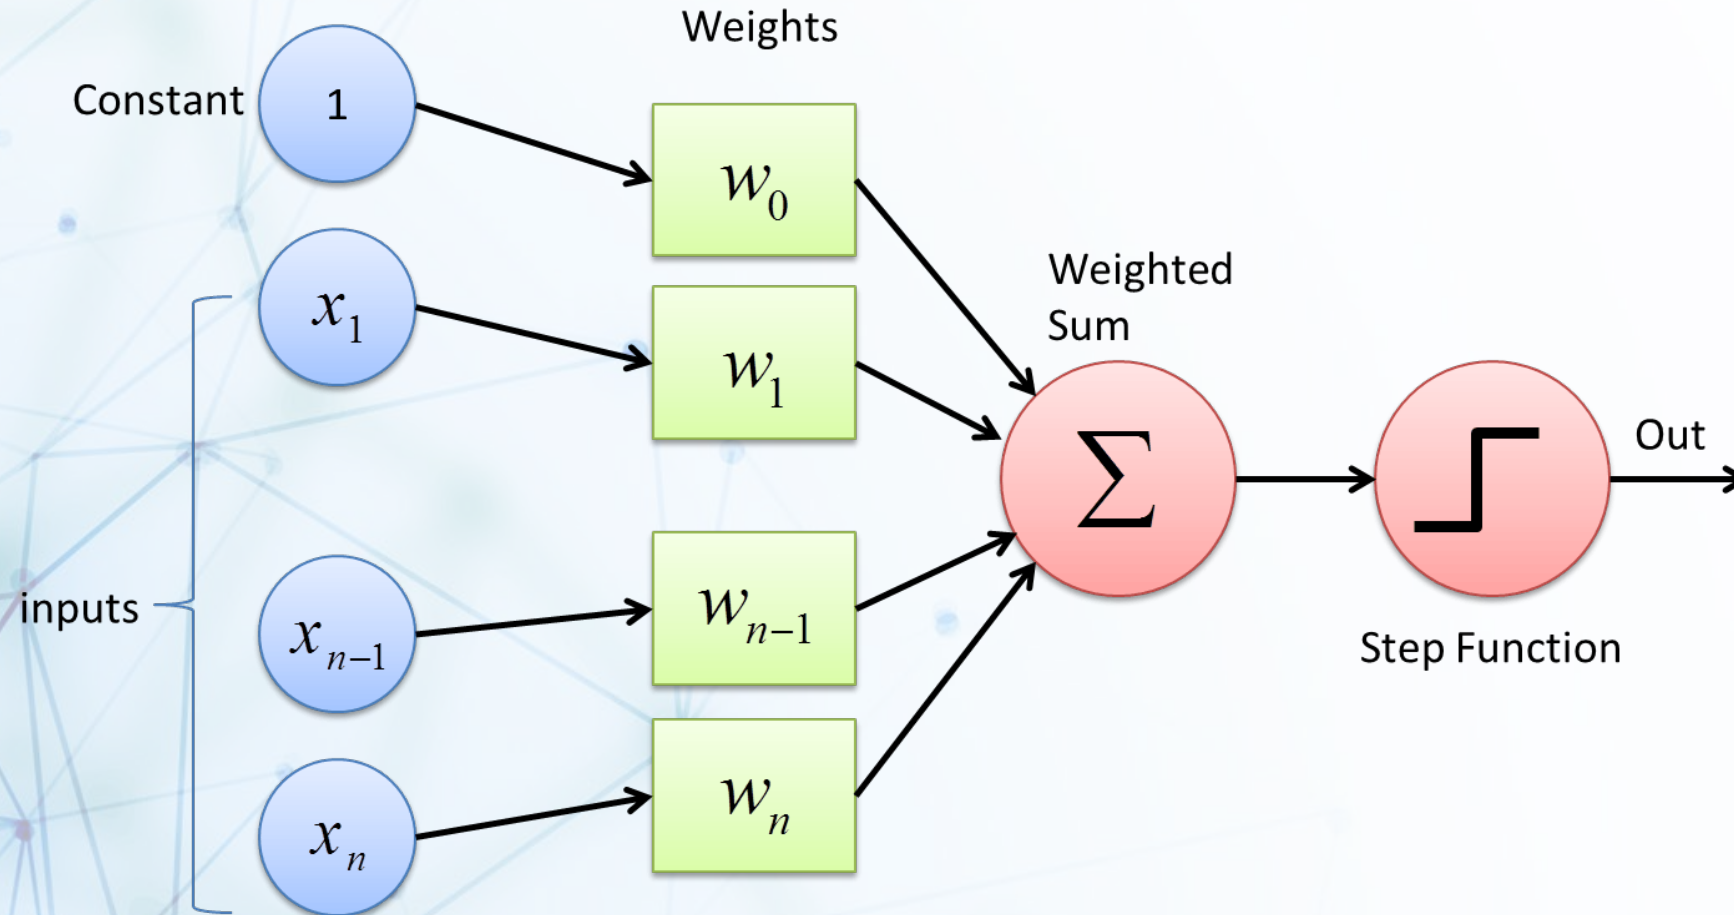
Neural networks



Source: wikipedia

Inspired by nature

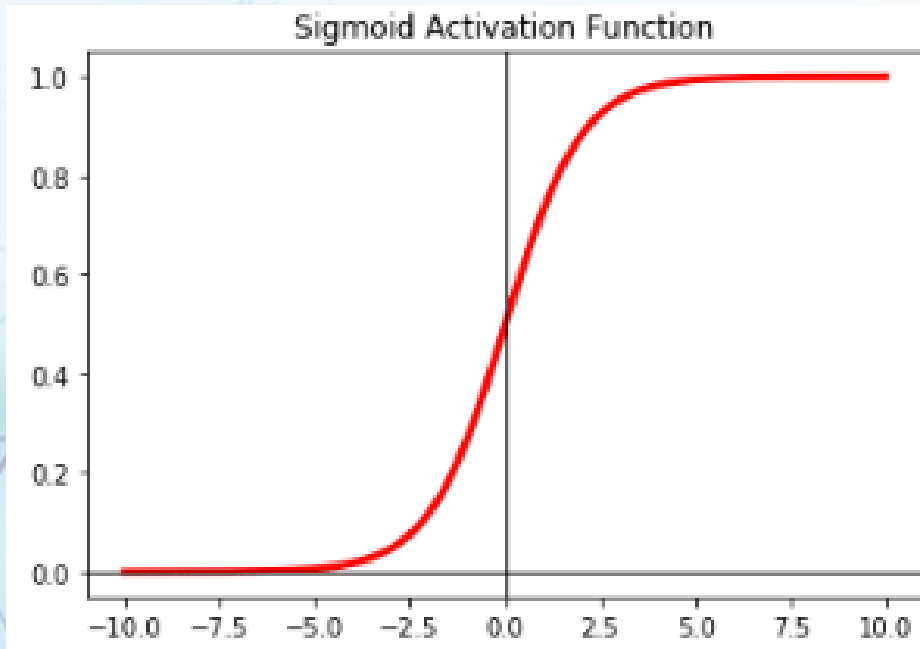
Artificial neurons or perceptrons




A perceptron is equal to linear or logistic regression

Activation functions

Sigmoid




Relu




Using non-linear activations any function can be approximated

Convolutional layers



Convolutions allow us to filter using less parameters and have symmetry!

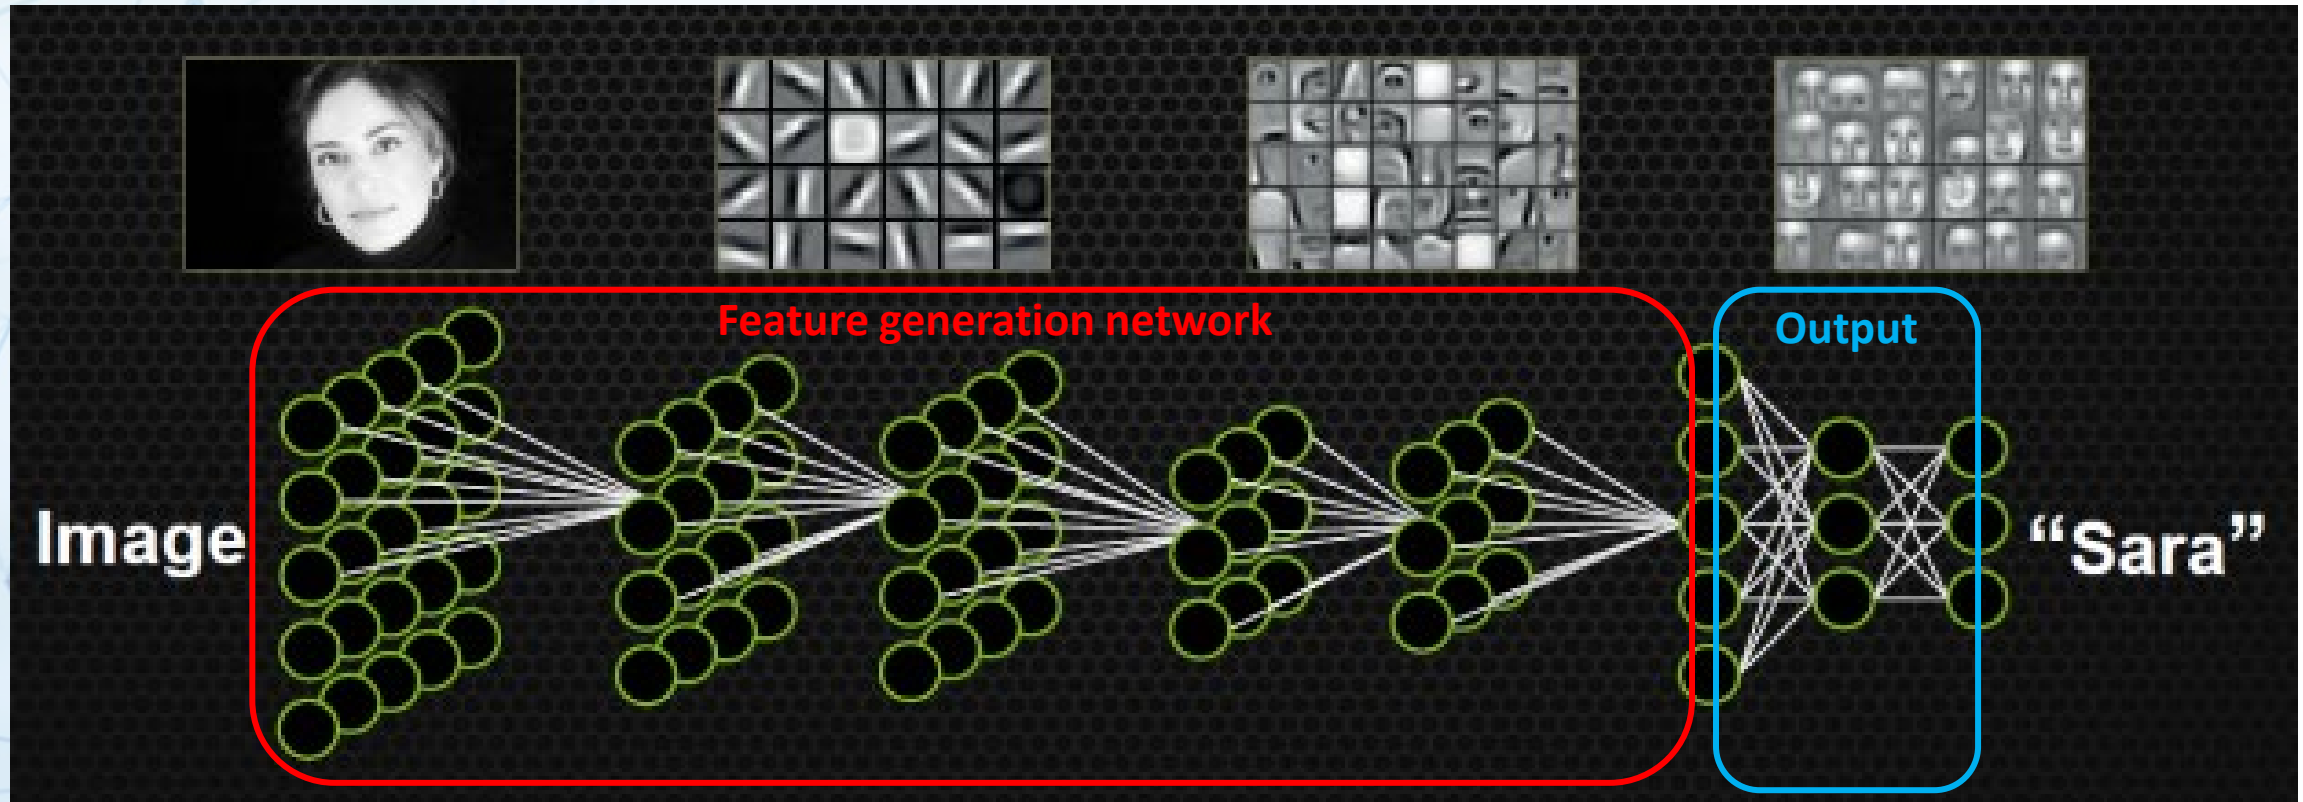
Deep learning



Source: NVIDIA Deep learning training

Features are engineered for you! but need lots of data... or do you?

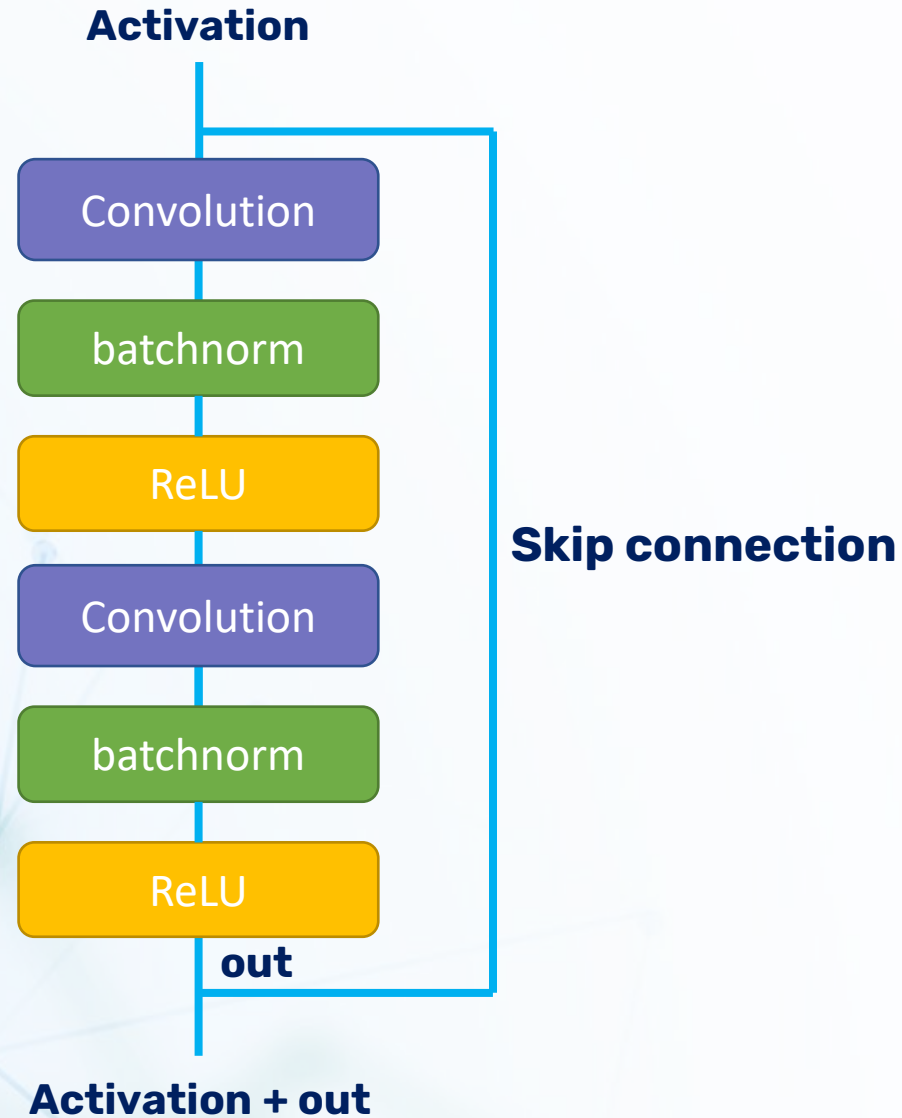
Transfer learning



Source: NVIDIA Deep learning training

With transfer learning we train on large datasets and finetune on small ones

Residual networks




Normalization and skip connections stabilize training

This session

- Analyze the images
- Choose a model
- Optimize the model
- Evaluate the results
- Interpret with explainable AI


Part of the pipeline needs to run in production

Training a neural network



Batches should contain enough variation

Augmentation



Adding some distortion to images improves our model's robustness

Loss functions

Regression

- MSE (L2)

$$\frac{\sum_N (target - pred)^2}{N}$$

- MAE (L1)

$$\frac{\sum_N |target - pred|}{N}$$

- Custom weights
- ...

Classification

- (Binary) Cross entropy,

$$-target * \log(prob\ pred) \ (target=1) \\ +(1 - target) * \log(1 - prob\ pred) \ (target = 0)$$


Summed over classes if multiclass

- For segmentation applied per pixel
- ...


The right metric guides the optimizer to the right goal

Loss surfaces

Normal net




Resnet



<https://proceedings.neurips.cc/paper/2018/file/a41b3bb3e6b050b6c9067c67f663b915-Paper.pdf>

Loss functions have complex surfaces with millions of parameters

Optimizing



3blue1brown - <https://mlfromscratch.com/optimizers-explained/#/>

Gradient descent allows us to stepwise optimize our parameters for our loss


More on optimizers

- <https://distill.pub/2017/momentum/> - try this yourself
- <https://runder.io/optimizing-gradient-descent/>
- <https://towardsdatascience.com/a-visual-explanation-of-gradient-descent-methods-momentum-adagrad-rmsprop-adam-f898b102325c>
- <https://towardsdatascience.com/understanding-backpropagation-algorithm-7bb3aa2f95fd>

Adding some distortion to images improves our model's robustness


Regularization

Dropout



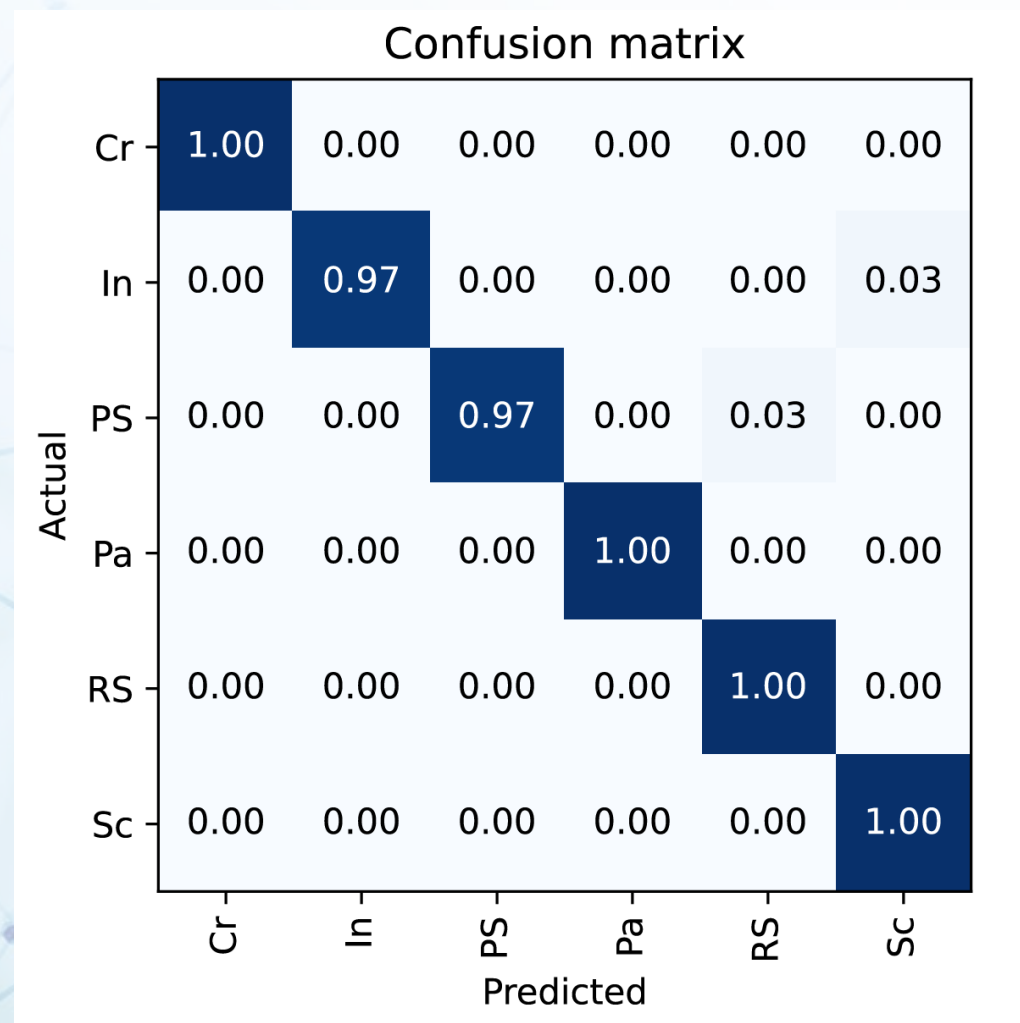
Randomly delete
neurons doing training

Weight decay



Add the norm of
weights to loss

Metrics




Metrics can, but don't have to be the same as a loss function (no backprop)

Top losses

Predicted	Actual	Loss	Probability
-----------	--------	------	-------------


Sc/In	/	2.82	/ 0.94
-------	---	------	--------



The best guess is not necessarily a good guess


Top losses

Softmax



Always gives the best prediction

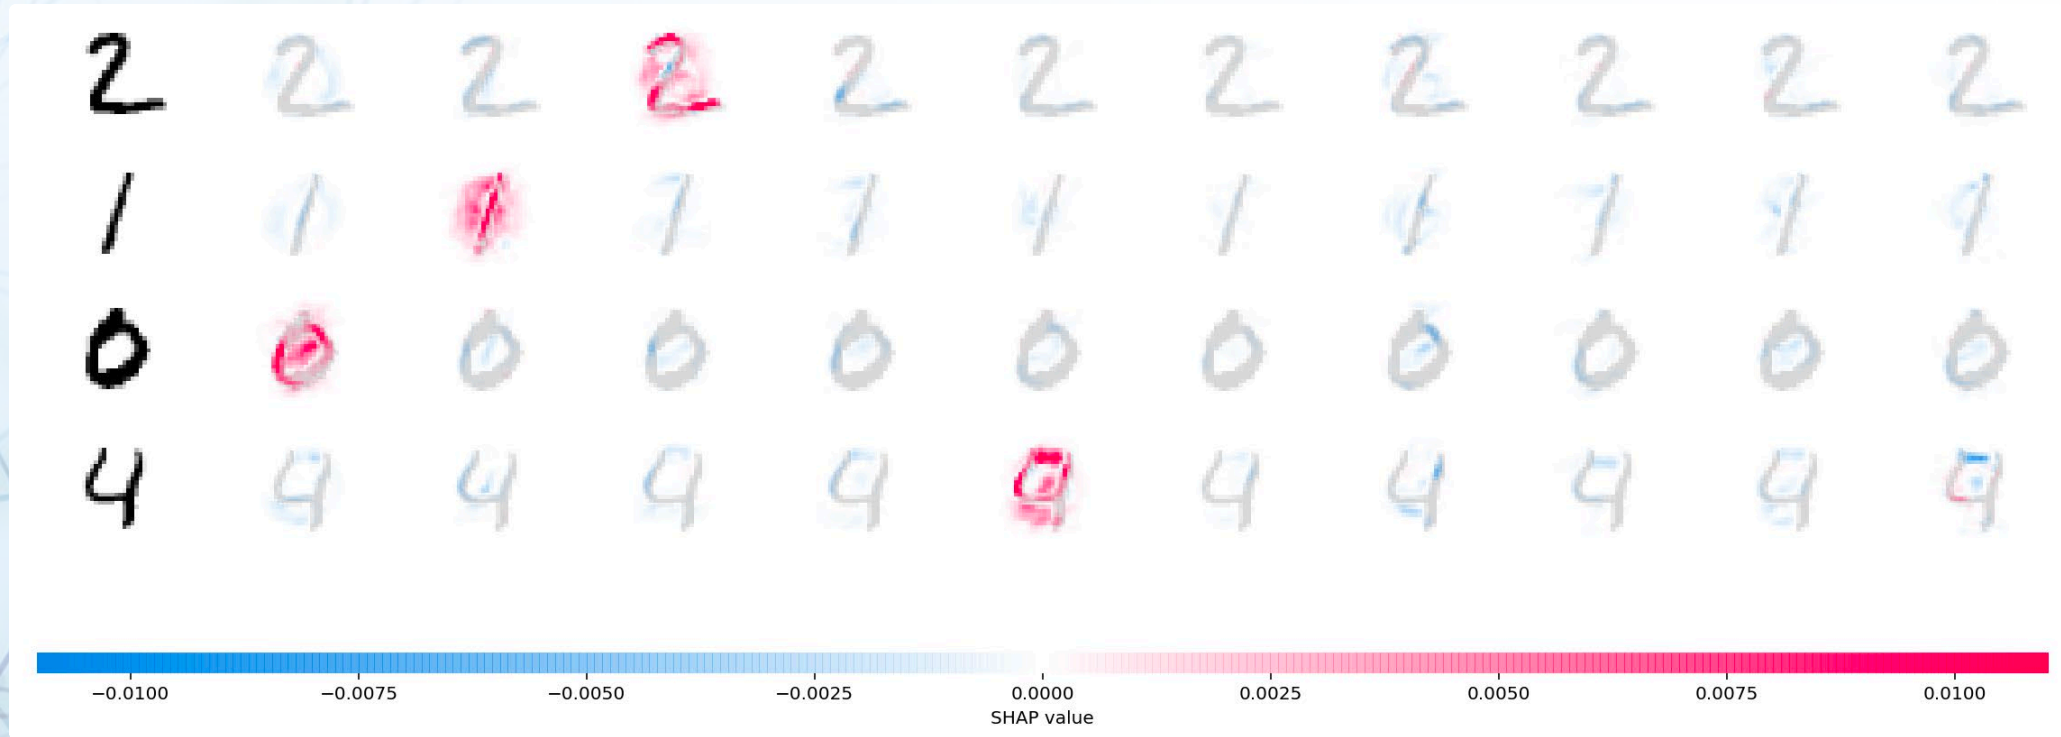
Sigmoid



Gives probability per class

SHAP


Our input features are pixels, can we trace them to the output?



<https://github.com/slundberg/shap>


Yes, by comparing to baseline images we can approximate shap values using grads

SHAP: scratches




Scratch is clearly highlighted

SHAP: large defects




Both dark and light regions used

SHAP: distributed defects



Distributed activation regions

Detecting problems with SHAP



Edges seem more important than they should be


INVESTEERT IN
JOUW TOEKOMST



View more online
<https://ai4mi.epotentia.com>

AI FOR **MATERIALS**
INDUSTRY

Glass



Materials discovery

Manufacturing



Sensor data

